

RL 2020/2021 Coursework Description

Due: March 31, 2021: 16:00

February 9, 2021

1 Introduction

The goal of this coursework is to implement different reinforcement learning algorithms covered in the lectures. By completing this coursework, you will get first-hand experience on how different algorithms perform in different decision-making problems.

Throughout this coursework, we will refer to lecture slides for your understanding and give page numbers to find more information in the RL textbook ("Reinforcement Learning: An Introduction (2nd edition)" by Sutton and Barto, <http://www.incompleteideas.net/book/RLbook2020.pdf>).

As stated in the course prerequisites, we do expect students to have a good understanding of Python programming, and of course any material covered in the lectures is the core foundation to work on this coursework. Many tutorials on Python can be found online.

We encourage you to start the coursework as early as possible to have sufficient time to ask any questions.

2 Contact

Piazza Please post questions about the coursework in the Piazza forum to allow everyone to view the answers in case they have similar questions. We provide different tags/folders in Piazza for each question in this coursework. Please post your questions using the appropriate tag to allow others to easily read through all the posts regarding a specific question.

Lab sessions There will also be demonstration sessions during which you can ask questions about the coursework. We highly recommend attending these sessions, especially if you have questions about PyTorch and the code base we use.

3 Getting Started

To get you started, we provide a repository of code to build upon. Each question specifies which sections of algorithms you are expected to implement and will point you to the respective files.

1. Installing Python3

The code base is fully written in Python and we expect you to use several standard machine learning packages to write your solutions with. Therefore, start by downloading Python to your local machine. We recommend you use at least Python version 3.6.

Python can be installed using the official installers (<https://www.python.org/downloads/>) or alternatively using a respective package-manager on Linux or Homebrew (<https://brew.sh>) on MacOS.

2. Create a virtual environment

After installing Python, we highly recommend creating a virtual environment to install the required packages. This allows you to neatly organise the required packages for different projects and avoid potential issues caused by insufficient access permissions on your machines. On Linux or MacOS machines, type the following command in your terminal:

```
python3 -m venv <environment name>
```

You should now see a new folder with the same name as the environment name you provided in the previous command. In your current directory, you can then execute the following command to activate your virtual environment on Linux or MacOS machines:

```
source <environment name>/bin/activate
```

If you are using Windows, please refer to the [official Python guide](#) for detailed instructions.

3. Download the code base to get started

Finally, execute the following command to download the code base:

```
git clone https://github.com/uae-agents/uae-rl2021.git
```

Navigate to **<Coursework directory with setup>** and execute the following command to install the code base and the required dependencies:

```
pip3 install -e .
```

For detailed instructions on Python's library manager `pip` and virtual environments, see the [official Python guide](#) and [this guide to Python's virtual environments](#).

4 Overview

The coursework contains a total of **100 marks** and counts towards **50% of the course grade**. Below you can find an overview of the coursework questions and their respective marks. More details on required algorithms, environments and required tasks can be found in Section 5. Submissions will be marked based on correctness and performance as specified for each question. Details on marking can be found in Section 6 and Section 7 presents instructions on how to submit your implementations and report correctly.

Question 1 – Dynamic Programming [15 Marks]

- Implement the following DP algorithms for MDPs
 - Value Iteration [7.5 Marks]
 - Policy Iteration [7.5 Marks]
-

Question 2 – Tabular Reinforcement Learning [20 Marks]

- Implement ϵ -greedy action selection [2 Marks]
 - Implement the following RL algorithms
 - Q-Learning [7 Marks]
 - On-policy first-visit Monte Carlo [7 Marks]
 - Tune the algorithms to solve Taxi-v3 and provide plot of average returns [4 Marks]
-

Question 3 – Deep Reinforcement Learning [35 Marks]

- Implement the following Deep RL algorithms
 - Deep Q-Networks [6 Marks]
 - REINFORCE [9 Marks]
 - Tune both algorithms to solve CartPole and provide plots of average returns [6 Marks]
 - Tune DQN to solve LunarLander and provide plot of average returns [6 Marks]
 - Provide a plot with analysis of the DQN losses during Cartpole training following the provided instructions [8 Marks]
-

Question 4 – Continuous Deep Reinforcement Learning [15 Marks]

- Implement DDPG for continuous RL [12 Marks]
 - Tune DDPG to solve Pendulum [3 Marks]
-

Question 5 – Multi-Agent Reinforcement Learning [15 Marks]

- Implement the following multi-agent RL algorithms to solve matrix games
 - Implement Independent Q-Learning [5 Marks]
 - Implement Joint-Action Learning with Opponent Modelling [10 Marks]

5 Questions

5.1 Question 1 – Dynamic Programming [15 Marks]

Description

The aim of this question is to provide you with better understanding of dynamic programming approaches to find optimal policies for Markov Decision Processes (MDPs). Specifically, you are required to implement the Policy Iteration (PI) and Value Iteration (VI) algorithms.

For this question, **you are only required to provide implementation of the necessary functions**. For each algorithm, you can find the functions that you need to implement under Tasks below. Make sure to carefully read the code documentation to understand the input and required outputs of these functions. We will mark your submission only based on the correctness of the outputs of these functions.

Algorithms

1. Policy Iteration (PI):

You can find more details including pseudocode in the RL textbook on page 80. Also see [Lecture 4 on dynamic programming](#) (pseudocode on slide 17).

2. Value Iteration (VI):

You can find more details including pseudocode in the RL textbook on page 83. Also see [Lecture 4 on dynamic programming](#) (pseudocode on slide 22).

Domain

In this exercise, we train dynamic programming algorithms on MDPs. We provide you with functionality which enables you to define your own MDPs for testing. For an example on how to use these functions, see the main function at the end of `exercise1/mdp_solver.py` where the MDP shown in Figure 1 is defined and given as input to the training function with $\alpha, \beta, R_{search}$, and R_{wait} set to 0.8, 0.4, 5, and 2 respectively.

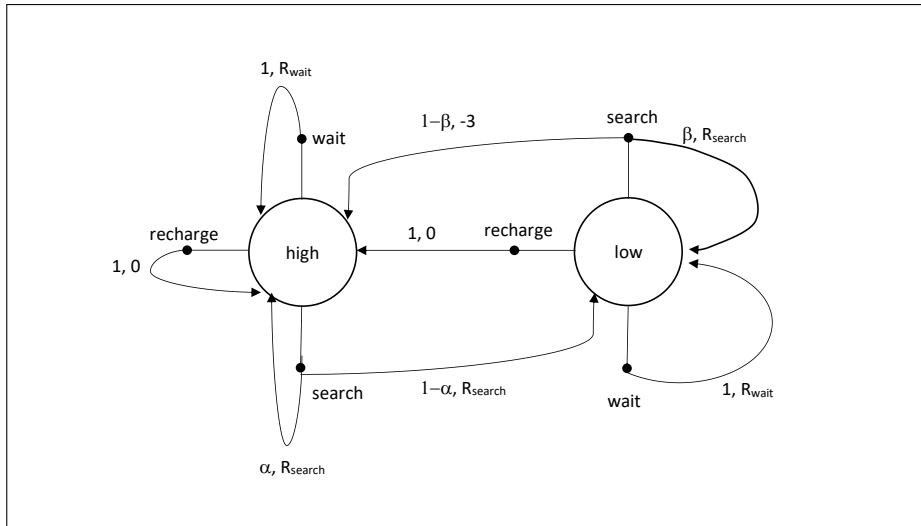


Figure 1: Example MDP for Exercise 1

As a side note, our interface for defining custom MDPs requires all actions to be valid over all states in the state space. Therefore, remember to include a probability distribution over next states for every possible state-action pair to avoid any errors from the interface.

Tasks

Use the code base provided in the directory `exercise1` and implement the following functions.

1. **Value Iteration**

[7.5 Marks]

To implement the Value Iteration algorithm, you must implement the following functions in the `ValueIteration` class:

- `_calc_value_func`, which must calculate the value function (table).
- `_calc_policy`, which must return the greedy deterministic policy given the calculated value function.

2. Policy Iteration

[7.5 Marks]

To implement the Policy Iteration algorithm, you must implement the following functions in the `PolicyIteration` class:

- `_policy_eval`, which must calculate the value function of the current policy.
- `_policy_improvement`, which must return an improved policy and terminate if the policy is stable.

Aside from the aforementioned functions, the rest of the code base for this question **must be left unchanged**. A good starting point for this question would be to read the code base and the documentations to get a better grasp how the entire training process works.

Directly run the file `mdp_solver.py` to print the calculated policies for VI and PI for a test MDP. Feel free to tweak or change the MDP and make sure it works consistently.

This question does not require a lot of effort to complete and you can provide a correct implementation with less than 50 lines of code. Additionally, training the method should require less than a minute of running time.

5.2 Question 2 – Tabular Reinforcement Learning [20 Marks]

Description

The aim of the second question is to provide you with practical experience on implementing model-free reinforcement learning algorithms with tabular Q-functions. Specifically, you are required to implement the **Q-Learning** (QL) and **on-policy first-visit Monte Carlo** (MC) algorithms.

For all algorithms, **you are required to provide implementations of the necessary functions**. You can find the functions that you need to implement below. Make sure to carefully read the documentation of these functions to understand their input and required outputs. We will mark your submission based on the **correctness of the outputs of the required functions** and the **performance of your learning agents measured by the average returns on the Taxi-v3 environment**.

Algorithms

1. Q-Learning (QL):
You can find more details including pseudocode for QL in the [RL textbook on page 131](#). Also see [Lecture 6 on Temporal Difference learning](#) (pseudocode on slide 19).
2. First-visit Monte Carlo (MC):
You can find more details including pseudocode for on-policy first-visit Monte Carlo with ϵ -soft policies in the [RL textbook on page 101](#). Also see [Lecture 5 on Monte Carlo methods](#) (pseudocode on slide 17).

Domain

In this question, we train agents on the OpenAI Gym Taxi-v3 environment. This environment is a simple task where the goal of the agent is to navigate a taxi (yellow box - empty taxi; green box - taxi with passenger) to a passenger (blue location), pick it up and drop it off at the destination (purple location) in a grid-world.




Figure 2: Rendering of two Taxi-v3 environment steps

The episode terminates once the passenger is dropped off at its destination or at a given maximum episode length (which can be set as a hyperparameter). The agent will be given a reward of -1 at each timestep, a reward of +20 for successfully delivering the passenger to its destination and -10 for executing actions pickup or dropoff illegally, i.e. trying to pickup a passenger at a location where no passenger is located or attempting to drop off without having a passenger in the taxi. Hence, the task consists of learning to navigate the grid-world and bringing the passenger as quickly to its target destination as possible.

A good hyperparameter scheduling for both algorithms should enable the agent to solve the Taxi-v3 environment. **We consider the environment to be solved when the agent can consistently achieve an average return of ≥ 7 and does not have “failed” episodes.** An episode is considered failed whenever the agent received negative returns in that particular episode.

Tasks

For this exercise, you are required to implement the functions listed below. Besides the correctness of these functions, we will also mark the performance achieved by your agents in the Taxi-v3 environment. See each paragraph below for more details on required functions, performance thresholds and respective marks.

Implementation

[16 Marks]

Use the code base provided in the directory `exercise2` and implement the following functions. All of the functions, that you need to implement for the three algorithms, are located in the `agents.py` file. Both algorithms to implement extend the `Agent` class provided in the script.

1. Base class

[2 Marks]

In the `Agent` class, implement the following function:

- `act`, where you must implement the ϵ -greedy exploration policy used by the Q-Learning and MC algorithm.

2. Q-Learning

[7 Marks]

To implement Q-Learning, you must implement the following functions in the `QLearningAgent` class:

- `learn`, where you must implement Q-value updates.
- `schedule_hyperparameters`, where you can schedule the values of Q-Learning hyperparameters to improve performance.

3. On-policy first-visit Monte Carlo

[7 Marks]

To implement the Monte Carlo with ϵ -soft policy algorithm, you must implement the following functions in the `MonteCarloAgent` class:

- `learn`, where you must implement the Monte Carlo with ϵ -soft policy Q-value update.
- `schedule_hyperparameters`, where you can schedule the values of the hyperparameters of MC to improve performance on the environment being used.

All other functions apart from the aforementioned ones **should not be changed**. All functions could be implemented with around 20 lines of code or less.

Performance

[4 Marks]

Besides correctness of the action selection and learning functions, we will also mark the performance of the Q-Learning and first-visit MC algorithms in the Taxi-v3 environment. You will receive the following marks given performance thresholds:

Performance marks	0/2	1/2	2/2
Q-Learning	< 0	< 7	≥ 7
First-visit MC	< 0	< 7	≥ 7

Table 1: Average (evaluation) returns required for given **performance marks** for Q-Learning and first-visit MC in the Taxi-v3 environment.

In addition to a correct implementation, hyperparameter tuning (adjusting hyperparameters in the `config` dictionaries in `train_q_learning.py` and `train_monte_carlo.py`) and scheduling (through `schedule_hyperparameters` functions) will be required to achieve full performance marks. Make sure that the performance you are getting is reliable by using the `plot_results.py` script and plotting evaluation returns across many seeds. You are free to modify this script if you need to plot something different but **you will need to submit your generated performance plots**.

Testing

You can find the training script for Q-Learning and Monte-Carlo on Taxi-v3 in `train_q_learning.py` and `train_monte_carlo.py` respectively. These execute training and evaluation using your implemented agents.

5.3 Question 3 – Deep Reinforcement Learning [35 Marks]

Description

In this question you are required to implement two Deep Reinforcement Learning algorithms: **DQN** [4] and **REINFORCE** [6] with function approximation.

In this task, you are **required to implement functions associated with the training process, action selection along with gradient-based updates done by each agent**. Aside from these functions, many components of the training process, along with the primary training setup have already been implemented in our code base. Below, you can find a list of functions that need to be implemented. Make sure to carefully read the documentation of functions you must implement to understand the inputs and required outputs of each component. We will mark your submission based on **the correctness of the functions you’ve implemented**, along with the **achieved average returns of the agents** on both specified domains using the hyperparameter configurations you’ve specified for training.

Algorithms

Before you start implementing your solutions, we recommend reading the original papers and looking at lectures and textbooks to provide you with better understanding of the details of both algorithms.

1. Deep Q-Networks (DQN):

DQN is one of the earliest Deep RL algorithms, which replaces the usual Q-table used in Q-Learning with a neural network to scale Q-Learning to problems with large or continuous state spaces. You can find more details including pseudocode for DQN in the [Nature publication](#) [4]. Also see [Lecture 12 on deep RL](#) (pseudocode on slide 17).

2. REINFORCE:

REINFORCE is an on-policy algorithm which learns a stochastic policy with gradient updates being derived by the policy gradient theorem (see Lecture 11, slide 11). You can find more details in the [publication](#) [6] and for pseudocode refer to Algorithm 1 provided below.

Domains

In this question, we train agents on the [OpenAI Gym CartPole](#) and [LunarLander](#) environments. CartPole is a well-known control task where the agent can move a cart left or right to balance a pole. The goal is to learn balancing the pole for as long as possible. Episodes are limited in length and terminate early whenever the pole tilts beyond a certain degree. The agent is rewarded for each timestep it achieves to maintain the pole in balance.

For the LunarLander task, the agent is in control of a small spaceship. The goal is to throttle engines to navigate the spaceship to land on a dedicated landing pad (marked by two flags). Rewards are assigned for controlled landing and the agent is punished for fuel consumption and crashing.




Figure 3: Rendering of the CartPole and LunarLander environments

A well-tuned implementation should achieve an average return higher than 195 on CartPole and 195 in the LunarLander environment.

Tasks

For this exercise, you are required to implement the functions listed below. Besides the correctness of these functions, we will also mark the performance achieved by your DQN and REINFORCE agents in the CartPole environment as well as the performance achieved by your DQN agent in the LunarLander environment. Furthermore, you are required to provide a brief PDF document visualising and explaining the loss of DQN during training on CartPole. See each paragraph below for more details on required functions, performance thresholds, the loss report and respective marks.

Implementation

[15 Marks]

Use the code base provided in the directory `exercise3` and implement the following functions. All of the functions which you need to implement for both algorithms are located in the `agents.py` file. Both algorithms to implement extend the `Agent` class provided in the script.

1. DQN

[6 Marks]

In `agents.py`, you will find the `DQN` class which you need to complete. For this class, implement the following functions:

- `__init__`, which creates a DQN agent. Here, you can set any hyperparameters and initialise any values for the class you need.
- `act`, which implements a ϵ -greedy action selection. Aside from the observation, this function also receives a boolean flag as input. When the value of this boolean flag is `True`, agents should follow the ϵ -greedy policy. Otherwise, agents should follow the greedy policy. This flag is useful when we interchange between training and evaluation.
- `update`, which receives a batch of experience from the replay buffer. Using experiences, which are tuples in the form of $\langle s_t, a_t, r_t, d_t, s_{t+1} \rangle$ gathered from the replay buffer, update the parameters of the value network to minimize:

$$\mathbb{L}_\theta = (r + \gamma(1 - d_t)\max_a Q(a|s_{t+1}; \theta') - Q(a_t|s_t; \theta))^2,$$

where θ and θ' are the parameters of the value and target network, respectively.

- `schedule_hyperparameters`, where you can implement a hyperparameter scheduling method that enables agents to perform well.

2. REINFORCE

[9 Marks]

The functions that you need to implement for REINFORCE are also located inside the `agents.py` file under the `Reinforce` class. For this class, provide the implementation of the following functions:

- `__init__`, which creates the REINFORCE agent. You can set additional hyperparameters and values required for training the agent here.
- `act`, which implements the action selection based on the stochastic policy produced by the policy network.
- `update`, which updates the policy based on the sequence of experience

$$\{\langle s_t, a_t, r_t, d_t, s_{t+1} \rangle\}_{t=1}^T$$

received by the agent during an episode. You must then implement a process that updates the policy parameters to minimize the following function:

$$\mathbf{L}_\theta = \frac{1}{T} \sum_{t=1}^T -\log(\pi(a_t|s_t; \theta))(G_t)$$

where θ are the parameters of the policy network, and G_t is the discounted reward-to-go calculated starting from timestep t .

- `schedule_hyperparameters`, where you can implement a hyperparameter scheduling method that enables agents to perform well.

You can find the pseudocode for REINFORCE below in Algorithm 1.

All other functions apart from the aforementioned ones **should not be changed**. In general, all of the required functions can be implemented with less than 20 lines of code.

Algorithm 1: REINFORCE: Monte-Carlo Policy Gradient

Output: $\pi(a|s, \theta^*)$: optimised parameterised policy**Input:** α : Learning rate γ : Discount factor**Initialise:** $\pi(a|s, \theta)$: Randomly initialise policy parameters θ

Loop forever (for each episode):Generate an episode $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$ following $\pi(\cdot|\cdot, \theta)$ $L_\theta \leftarrow 0$ // Initialise loss to 0 $G \leftarrow 0$ // Initialise the returns to 0Loop backward in the episode $t = T - 1, \dots, 0$: $G \leftarrow R_{t+1} + \gamma G$ $L_\theta \leftarrow L_\theta - G \log \pi(A_t|S_t, \theta)$ $L_\theta \leftarrow L_\theta / T$ Perform a gradient step with learning rate α on L_θ with respect to θ

Performance**[12 Marks]**

Besides correctness of the aforementioned algorithms, we will also mark the performance of DQN and REINFORCE in the CartPole environment and DQN performance in the LunarLander environment. See Table 2 for performance marks received for specified evaluation return thresholds. Note, that you are required to tune both DQN and REINFORCE on the CartPole environment, but we only mark performance of DQN in the LunarLander environment. In particular, **we do not mark performance of REINFORCE in the LunarLander environment.**

Performance marks	0/3	2/3	3/3
DQN	< 150	< 195	≥ 195
REINFORCE	< 150	< 195	≥ 195

(a) Average (evaluation) returns required for given **performance marks** for CartPole.

Performance marks	0/6	2/6	4/6	6/6
DQN	< 100	< 150	< 195	≥ 195

(b) Average (evaluation) returns required for given **performance marks** for LunarLander.

Table 2: Average (evaluation) returns required for given **performance marks** for (a) DQN and REINFORCE on the CartPole environment and for (b) DQN on the LunarLander environment.

In order to reliably solve CartPole and LunarLander, tune the hyperparameters of DQN and REINFORCE. To do so, switch to the respective configuration and change the hyperparameters as you see fit (either through the config dictionaries themselves, or the `schedule_hyperparameters` functions). Make sure that the performance you are getting is reliable by using the `plot_results.py` script and plotting many seeds. You are free to modify this script if you need to plot something different but **you will need to submit your generated performance plots**. Also, **you will need to provide us with saved parameters for DQN in LunarLander so that we can verify the performance**. We will not train your DQN agent on LunarLander, but just restore the parameters you provide and evaluate.

Understanding the Loss**[8 Marks]**

This part of the exercise will attempt to further your understanding of the loss function in DQN. First, plot the DQN loss with “timesteps trained” on the x-axis and the loss itself in the y-axis. The values should be generated on a single run of CartPole. Note, we already provide you with the functionality to collect and plot the DQN loss as required. Simply set the “`plot_loss`” value within the `CARTPOLE_CONFIG` in `train_dqn.py` to `True` and you should receive a plot as stated at the end of training.

In machine learning, it is often expected for the value of the loss to drop during training. Is it

also happening in your plot? Create and submit a separate `loss.pdf` document, which ...

- includes your plot of the DQN loss as specified and
- in **less than 10 lines** (using a fontsize of 12pt)
 - i) explain why the loss is not behaving as in typical supervised learning approaches (where we usually see a fairly steady decrease of the loss throughout training) and
 - ii) provide an explanation for the spikes which are standing out in the plot.

Testing

To test your implementation, we provide you with two scripts which execute your DQN and REINFORCE implementations. You can find the scripts inside `train_dqn.py` and `train_reinforce.py` to train DQN and REINFORCE, respectively. Inside these scripts, we provide you with configurations that enable you to train the methods in the CartPole and LunarLander environments. To better understand how your implemented functions are used in the training process, read the code and documentation provided in these scripts.

Change hyperparameters specified inside the configurations to identify well performing agents. If your implementation is correct and hyperparameters are well configured, you should be able to achieve given returns of 195 in CartPole and LunarLander. Due to its simplicity, we highly recommend you to try training on the CartPole environment first. For a correct implementation, the training process requires less than 2 minutes to train CartPole and about 20 minutes to train the LunarLander agent.

Make sure that you save the model parameters for LunarLander using the `save` function included. Make sure the parameters load correctly when calling the `restore` function. **Include the saved parameters in your submission.**

We provide you with some hyperparameters that should work adequately out-of-the-box. However, the performance, measured by average returns, of your agents will be tested using your given configuration and as such, tuning the hyperparameters could improve the performance and is highly recommended.

5.4 Question 4 – Continuous Deep Reinforcement learning [15 Marks]

Description

So far, we implemented algorithms such as DQN and REINFORCE which define value functions and policies, respectively, for discrete actions, i.e. each action in a state is assigned a specific value or action selection probability. However, in some problems such as control in robotics there might be continuous actions e.g. representing force which is applied by a motor. To be able to learn policies for such continuous action spaces, we need different RL techniques. The goal of this question is to provide you with experience on (deep) RL algorithms which can be applied in such continuous action spaces. To achieve this aim, you are required to implement the **Deep Deterministic Policy Gradient** (DDPG) [3] algorithm and train it to solve a Pendulum control task.

Algorithm

Deep Deterministic Policy Gradient (DDPG) [3] is building on top of Deterministic Policy Gradient (DPG) [5] and extending this RL algorithm for continuous action spaces with function approximators. We highly recommend reading the DDPG paper in addition to lecture materials to familiarise yourself with the algorithm. In contrast to discrete action environments, where an action is a scalar integer, the action in continuous action environments is an N-dimensional vector where, N is the dimension of the action space. Therefore, the Q-network in DDPG outputs a value estimate given a state and action in contrast to just receiving a state in DQN. Additionally, the action space usually has an upper and a lower bound.

For example, imagine a car with two-dimensional action space, throttle and turn, where throttle takes values in $[-1, 1]$, and turn takes values in $[-45, 45]$. At each timestep, the controlled agent should return a two-dimensional action, where the first element represents the throttle and should be in the range of $[-1, 1]$, and the second element represents the turn and therefore should be in the range of $[-45, 45]$. In the environment used for this question, all actions are bounded by $[-2, 2]$. The actor network is initialised with a tanh activation function and scaled with a factor of 2 to compute actions in the range $[-2, 2]$. You **must not** change this initialisation.

Please note that an epsilon-greedy policy, which was applied in DQN, cannot be applied in continuous action environments, because the number of possible actions are infinite. For this reason, we use Gaussian noise \mathcal{N} to explore.

$$a = \mu(s) + \eta$$

$$\eta \sim \mathcal{N}(\mathbf{m}, \sigma)$$

For this exercise, we consider that the noise is a Gaussian function with mean $\mathbf{m} = \mathbf{0}$ and standard deviation $\sigma = 0.1\mathbf{I}$ for identity matrix \mathbf{I} .

Using experiences, which are tuples in the form of $\langle s_t, a_t, r_t, d_t, s_{t+1} \rangle$ gathered from the replay buffer, update the parameters of the critic network to minimize:

$$\mathbb{L}_\theta = (r + \gamma(1 - d_t)Q(\mu(s_{t+1}; \phi'), s_{t+1}; \theta') - Q(a_t, s_t; \theta))^2,$$

where θ and θ' are the parameters of the critic and target critic network, respectively, and ϕ' are the parameters of the target actor network. Using the same batch, implement and minimise the deterministic policy gradient error to update the parameters of the actor:

$$\mathbb{L}_\phi = -Q(\mu(s_t; \phi), s_t; \theta)$$

where ϕ are the parameters of the actor's network. The gradient flows through the critic network back to the parameters of the actor. Please note, that during the update of the actor's parameters, the parameters of the critic network should remain fixed and not be updated.

Domain

In this question, we ask you to train agents in the OpenAI Gym Pendulum environment. Pendulum is a control task where an agent can apply force to balance a pendulum upwards. The goal is to learn to bring and keep the pendulum in an upward position. The agent observes the angle of the pendulum and chooses an action representing the torque applied to the pendulum. The agent is rewarded for keeping the pendulum in an upward position. A well-tuned implementation should achieve an average return higher than -300 .



Figure 4: Rendering of two Pendulum environment steps

Tasks

For this exercise, you are required to implement the functions listed below. Besides the correctness of these DDPG functions, we will also mark the performance achieved by your DDPG agent in the Pendulum environment. See each paragraph below for more details on required functions, performance thresholds and respective marks.

Implementation

[12 Marks]

Use the code base provided in the directory `exercise4` and implement the following functions. In `agents.py`, you will find the DDPG class which you need to complete. For this class, implement the following functions:

- `__init__`, which creates a DDPG agent. Here, you have to initialise the Gaussian noise. Use the imported class from `torch.distributions`, `Normal`, to define a noise variable. During exploration you should call the function `sample()` from the `Normal` instance. Also, you can set any additional hyperparameters and initialise any values for the class you need.
- `act`, which implements the action selection method of DDPG. Aside from the observation, this function also receives a boolean flag as input. When the value of this boolean flag is `True`, agents should follow an exploratory policy using noise as specified above. Otherwise, agents should follow the deterministic policy without any noise. This flag is useful when we interchange between training and evaluation.

Hint: Remember to clip the action in the range of $[-2, 2]$ before returning the action.

- `update`, which receives a batch of experience from the replay buffer. Using experiences, which are tuples in the form of $\langle s_t, a_t, r_t, d_t, s_{t+1} \rangle$ gathered from the replay buffer, update the parameters of the critic network to minimize:

$$\mathbb{L}_\theta = (r + \gamma(1 - d_t)Q(\mu(s_{t+1}; \phi'), s_{t+1}; \theta') - Q(a_t, s_t; \theta))^2,$$

where θ and θ' are the parameters of the critic and target critic network respectively, and ϕ' are the parameters of the target actor network. Using the same batch implement and minimise the deterministic policy gradient error to update the parameters of the actor:

$$\mathbb{L}_\phi = -Q(\mu(s_t; \phi), s_t; \theta)$$

where ϕ are the parameters of the actor's network. The gradient flows through the critic network back to the parameters of the actor. Please note, that during the update of the actor's parameters, the parameters of the critic network should remain fixed and not be updated.

- `schedule_hyperparameters`, where you can implement a hyperparameter scheduling method that enables agents to perform well.

Performance

[3 Marks]

Besides correctness of the action selection and learning functions, we will also mark the performance of the DDPG algorithm in the Pendulum environment. You will receive the following marks given performance thresholds shown in Table 3.

In addition to a correct implementation, hyperparameter tuning (adjusting hyperparameters in the `config` in `train_ddpg.py`) and scheduling (through `schedule_hyperparameters`) will be

Performance marks	0/3	1/3	2/3	3/3
DDPG	< -500	< -400	< -300	≥ -300

Table 3: Average (evaluation) returns required for given **performance marks** for DDPG in the Pendulum environment.

required to achieve full performance marks. Also, you will need to provide us with saved parameters for DDPG in Pendulum so that we can verify the performance. Make sure that the performance achieved by your saved parameters (saved at the end of training in `train_ddpg.py`) are reliable by using the `evaluate_ddpg.py` script over many evaluation episodes.

5.5 Question 5 – Multi-Agent Reinforcement Learning [15 Marks]

Description

The aim of this question is to provide you with experience on Multi-Agent Reinforcement Learning (MARL). To achieve this aim, you are required to implement two tabular MARL algorithms **Independent Q-Learning (IQL)** and **Joint Action Learning with Opponent Modelling (JAL-OM)** [1] to solve a simple version of the **penalty game** and attempt to solve the **climbing game** [2]. You can find more information on both algorithms in Lecture 15 on Multi-agent learning II.

For this question, you are required to submit your code that implements the **IQL** and **JAL-OM** algorithms.

Algorithms

1. Independent Q-Learning (IQL):

Independent learning as a general concept (see Lecture 15 on Multi-agent learning II) is one of the simplest approaches to MARL. Essentially, each agent is following a single-agent RL algorithm and ignores the presence of other agents by considering them part of the environment. In IQL, each agent follows the (tabular) Q-Learning algorithm. We recommend taking a look at pseudocode and instructions of Q-Learning, see Section 5.2.

2. Joint Action Learning with Opponent Modelling (JAL-OM):

In joint action learning, agents consider the presence and action selection of other agents. For this particular algorithm, simple opponent modelling is used to inform the value function. Each agent maintains a model aiming to approximate the policy of the other agent(s). You can find more details for JAL-OM in Lecture 15 on multi-agent learning II. For pseudocode and all intended details, please refer to Algorithm 2 below:

Algorithm 2: Joint Action Learning with Opponent Modelling (controlling agent i)

Output:

$s, \mathbf{a} : Q_i(s, \mathbf{a})$: Q-values for all state-(joint-)action pairs

Input:

α : Learning rate
 ϵ : Epsilon for *epsilon*-greedy policy
 γ : Discount factor
 T : Maximum training steps

Initialise:

$\forall s, \mathbf{a} : Q_i(s, \mathbf{a}) = 0$: Initialise Q-values,
 $\forall s : N(s) = 0$: Initialise state (visitation) counts,
 $s, \mathbf{a}_{-i} : C(s, \mathbf{a}_{-i}) = 0$: Initialise state-action pair counts for actions of all other agents (but i)

Observe the initial state of the environment $s \leftarrow s_0$.

for t in $0, \dots, T - 1$ **do**

if Act random (with probability ϵ) // Action Selection
 then
 Choose random action a_i
 else
 Choose best-response action $a_i \leftarrow \arg \max_{a_i} EV_i(s, a_i)$
 end
 Observe next state s' and own reward r_i after applying joint action $\mathbf{a} = (a_1, \dots, a_N)$
 $N(s) \leftarrow N(s) + 1$ // Update Visitation Count
 $C(s, \mathbf{a}_{-i}) \leftarrow C(s, \mathbf{a}_{-i}) + 1$ // Update Model(s)
 $Q_i(s, \mathbf{a}) \leftarrow Q_i(s, \mathbf{a}) + \alpha [r_i + \gamma \max_{a'_i} EV_i(s', a'_i) - Q_i(s, \mathbf{a})]$ // Update Q-values

end

with

$$EV_i(s, a_i) = \sum_{\mathbf{a}_{-i}} \frac{C(s, \mathbf{a}_{-i})}{N(s)} Q(s, (a_i; \mathbf{a}_{-i}))$$

Domain

In this question, we ask you to train agents on two matrix games, the **penalty game** and the **climbing game**. Details on both games can be found in [2]. The payoff matrices of the climbing game and the penalty game, respectively, are:

	a0	a1	a2
b0	0	6	5
b1	-30	7	0
b2	11	-30	0

(a)

	a0	a1	a2
b0	k	0	10
b1	0	2	0
b2	10	0	k

(b)

Table 4: Payoff Matrix for (a) climbing game and (b) penalty game with penalty k .

Since these environments are matrix games, there are no states defined for these games. To enable you to use the algorithms, which are designed for sequential decision making problems, we emulate the interaction in matrix games by always returning a zero constant as the state, and by setting the done flag to **True** each time the agents provide actions to the environment. Hence, each "episode" will just consist of a single interaction. For this exercise, we use a penalty $k = -5$ for the penalty game. A good hyperparameter scheduling for both algorithms should enable the agents to solve the penalty game optimally (reaching payoff of 10 for each agent) and reach the Nash Equilibrium of (a_1, b_1) with payoff 7 for the climbing game.

Tasks

Use the code base provided in the directory **exercise5** and implement the following functions. All of the functions that you need to implement for the two algorithms are located in the **agents.py** file. Both algorithms to implement extend the **MultiAgent** class provided in the script.

1. Independent Q-Learning [5 Marks]

To implement Independent Q-Learning, you must implement the following functions in the **IndependentQLearningAgents** class:

- **act**, where you must implement the action selection for all agents. Each agent chooses an action independently following a ϵ -greedy policy given its Q-table.
- **learn**, where you must implement the updates to each agents' Q-table.
- **schedule_hyperparameters**, where you can schedule the values of the hyperparameters of Independent Q-Learning to improve performance on the environments being used.

2. Joint Action Learning with Opponent Modelling [10 Marks]

To implement Joint Action Learning with Opponent Modelling, you must implement the following functions in the **JointActionLearning** class:

- **act**, where you must implement the action selection for all agents. Each agent chooses an action independently following a ϵ -greedy policy given its expected values for its actions in the current state.
- **learn**, where you must implement the Joint Action Learning updates to each agents' Q-table and models.
- **schedule_hyperparameters**, where you can schedule the values of the hyperparameters of JAL-OM to improve performance on the environments being used.

You can find the pseudocode for JAL-OM above in Algorithm 2.

All other functions apart from the aforementioned ones **should not be changed**. All functions could be implemented with around 40 lines of code or less.

Testing

You can find the training script for IQL and JAL-OM for both matrix games in **train_iql.py** and **train_jal.py**, respectively. These execute training and evaluation using your implemented agents.

6 Marking

Academic Conduct Please note that any assessed work is subject to University regulations and students are expected to follow any such regulations on academic conduct:

<http://web.inf.ed.ac.uk/infweb/admin/policies/academic-misconduct>

Correctness Marking As mentioned for most questions, we partly mark your submissions based on the correctness of the implemented functions. For pre-defined functions we ask you to implement, including most functions stated across all questions, we use unit testing scripts. In these scripts, we pass the same input into both your and our reference implementation and assign you marks according to whether the output of your function matches the expected output provided by our reference implementation. For functions which are evaluated for correctness, you must read the documentation to ensure that your implementation follows the expected format. **Only change files and functions specified for Questions 1–5! Otherwise, automated marking might fail which could lead to a deduction in marks.**

Performance Marking For most performance evaluation, we will run the training scripts of the code base to ensure that your agent solves the environments we used for training measured by the achieved average returns. Additionally, provide the plots as instructed since this is also part of performance evaluation. You will be provided with marks depending on the achieved average returns, see the respective tables in each exercise with performance marks for details. Therefore, **make sure that the hyperparameters of your algorithms have been appropriately tuned and are set in the configurations of the respective training scripts to achieve the required thresholds.** Also, for Question 3 and 4, **make sure to provide saved model parameters for DQN on LunarLander and DDPG trained on Pendulum** as instructed in the respective questions.

Writeup Marking For Question 3, do not forget to include a write-up required for the last task about DQN loss during Cartpole training. This document is expected to be a **PDF document** using a **fontsize of 12pt**. We will mark the write-up based on the quality of explanation about the DQN loss development and spikes as well as the plot itself. **Make sure to provide a clear and concise answer since you are only allowed to submit up to 10 lines for the write-up.**

7 Submission Instructions

Before you submit your implementations, make sure that you have organised your files according to the structure indicated in Figure 5. Your submission should have the same structure as the code base we’ve provided for this coursework with the addition of files we required you to submit, indicated by the bold font in the Figure. For other scripts we’ve provided in the code base, make sure you have implemented all the required functions.

Finally, compress the **r12021** folder into a **zip** file and submit the compressed file through Learn. In your Learn page, you can choose the **Coursework and Exam** panel and find the **Coursework 1 Submission** page. For general guidance on submitting files through Learn, you can find further information through the blog post linked below:

<https://blogs.ed.ac.uk/ilts/2019/09/27/assignment-hand-ins-for-learn-guidance-for-students/>

Late Submissions All submissions are timestamped automatically and **we will mark the latest submission**. If you submit your work after the deadline a late penalty will be applied to this submission unless you have received an approved extension. Please be aware that marking for late submissions may be delayed and marks may not be returned within the same timeframe as for on-time submissions.

For additional information or any queries regarding late penalties and extension requests, follow the instructions stated on the School web page below:

web.inf.ed.ac.uk/infweb/student-services/ito/admin/coursework-projects/late-coursework-extension-requests

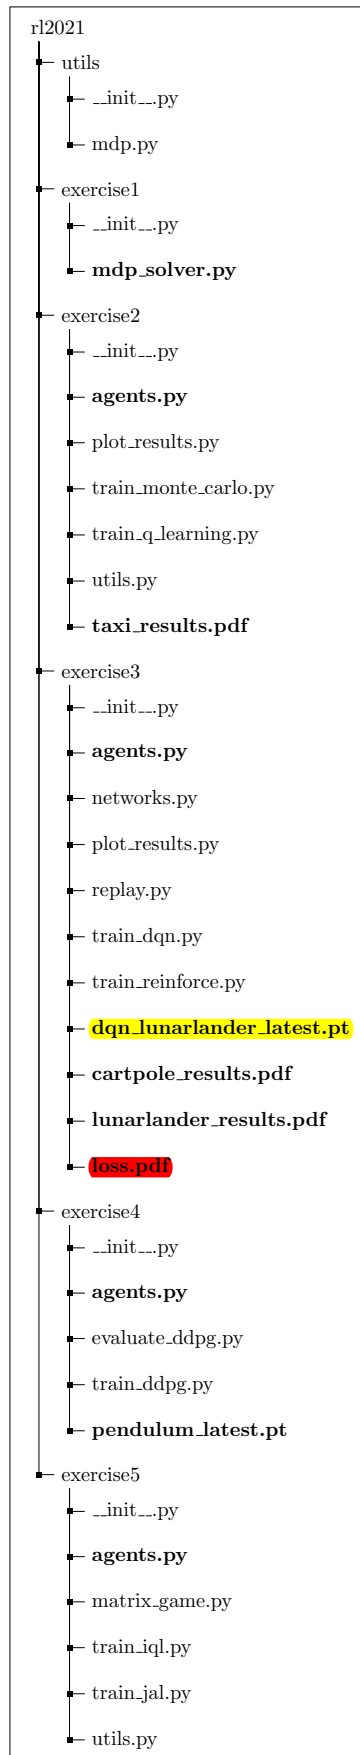


Figure 5: Required folder structure for submission. Files which need to be modified or created for this coursework are marked in **bold**.

References

- [1] Michael Bowling and Manuela Veloso. “Multiagent learning using a variable learning rate”. In: *Artificial Intelligence* 136.2 (2002), pp. 215–250.
- [2] Caroline Claus and Craig Boutilier. “The dynamics of reinforcement learning in cooperative multiagent systems”. In: *AAAI/IAAI* 1998.746-752 (1998), p. 2.
- [3] Timothy P Lillicrap et al. “Continuous control with deep reinforcement learning”. In: *International Conference on Learning Representations* (2015).
- [4] Volodymyr Mnih et al. “Human-level control through deep reinforcement learning”. In: *Nature* 518.7540 (2015), pp. 529–533.
- [5] David Silver et al. “Deterministic policy gradient algorithms”. In: 2014.
- [6] Richard S Sutton et al. “Policy gradient methods for reinforcement learning with function approximation”. In: *Advances in Neural Information Processing Systems*. 2000, pp. 1057–1063.